

Current AI systems may be highly capable while still making decisions within a narrow evaluation horizon.

This project tests whether agents become safer and more capable when they explicitly reason beyond task completion—maintaining uncertainty, checking irreversibility, and preserving future options before committing to actions that cannot be undone.

$$NEXT_i(t) > 0$$

= system **i** still has at least one possible continuation

$$NEXT_i(t) > 0 \wedge NEXT_i(t + 1|a) = 0 \wedge \neg Recover_i(a)$$

⇒ deeper deliberation before commitment

Fast reasoning → detect uncertainty / irreversibility → deeper deliberation